

# Swabha Swayamdipta

GABILAN ASSISTANT PROFESSOR OF COMPUTER SCIENCE, UNIVERSITY OF SOUTHERN CALIFORNIA

✉ swabhas@usc.edu | 🏠 <https://swabhs.com> | 🌐 swabhs | 🐦 @swabhaz

## Research Interests

NATURAL LANGUAGE PROCESSING AND MACHINE LEARNING

- Generative Evaluation: Evaluation of generated language via distributional effects or along specific axes
- Data Interpretability: Estimating the usefulness of data sources for model training and evaluation
- Model (Mechanistic) Interpretability: Understanding predictive behavior and properties of models that uniquely identify them
- Language Generation: Improved theory behind inference and algorithms for inference with and without linguistic constraints
- Language and Society: Language technologies to study societal biases in humans and models

## Professional Experience

### University of Southern California

GABILAN ASSISTANT PROFESSOR OF COMPUTER SCIENCE

- Associate Director of the Center for AI and Society, since Fall 2023

Los Angeles, CA, USA

Aug 2022 - PRESENT

### Allen Institute for AI

POSTDOCTORAL INVESTIGATOR

- **Advisor:** Yejin Choi
- **Project Team:** MOSAIC

Seattle, WA, USA

Aug 2019 - June 2022

### Oracle Server Technologies

MEMBER TECHNICAL STAFF

- **Project:** Unified messaging for Fusion Middleware

Bangalore, India

2010 - 2011

## Education

### Carnegie Mellon University

PHD IN LANGUAGE AND INFORMATION TECHNOLOGIES

- **Advisors:** Noah A. Smith, Chris Dyer
- **University of Washington**, Seattle, WA — Visiting PhD student Fall 2015 - Summer 2019
- **Thesis:** Syntactic Inductive Biases in NLP

Pittsburgh, PA, USA

2013 - 2019

### Columbia University

MASTERS IN COMPUTER SCIENCE

- **Advisors:** Owen Rambow, Michael Collins
- **Specialization:** Natural Language Processing

City of New York, NY, USA

2011 - 2012

### National Institute of Technology

BACHELORS OF TECHNOLOGY IN COMPUTER SCIENCE AND ENGINEERING

Calicut, Kerala, India

2006 - 2010

## Awards and Honors

- |      |   |                                    |
|------|---|------------------------------------|
| 2023 | <b>Rising Stars Award</b> , Intel   |                                    |
| 2022 | <b>Young Investigators Award</b> , Allen Institute for AI   |                                    |
| 2022 | <b>Gabilan Assistant Professor Fellowship</b> , Women in Science and Engineering  | USC                                |
| 2022 | <b>Outstanding Paper Award</b> , Understanding Dataset Difficulty with V-Usable Information                             | ICML                               |
| 2021 | <b>Outstanding Paper Award</b> , MAUVE: Measuring the Gap Between Neural Text and Human Text using Divergence Frontiers | NeurIPS                            |
| 2020 | <b>Best Paper: Honorable Mention</b> , Don't Stop Pretraining: Adapt Language Models to Domains and Tasks               | ACL                                |
| 2014 | <b>Best Student Presentation</b> , Diversity in Dependency Parsing  | CMU LTI Student Research Symposium |
| 2009 | <b>Sun Campus Ambassador Scholarship</b> , Bangalore, India   | Sun Microsystems                   |
| 2003 | <b>National Talent Search Scholarship</b> , Rourkela, India   | NCERT, India                       |

## Mentorship

Spring 2024-Present	<b>Brihi Joshi</b> , PhD, USC	
Fall 2023-Present	<b>Matthew Finlayson</b> , PhD, USC	
Fall 2022-Present	<b>Sayan Ghosh</b> , PhD, USC	
Fall 2023-Present	<b>Jaspreet Ranjit</b> , PhD, USC	
Spring 2024-Present	<b>Shauryasikt Jena</b> , Masters, USC	
Spring 2024-Present	<b>Xingjian Dong</b> , Masters, USC	
Spring 2023-Present	<b>Xinyue Cui</b> , Masters, USC	
Summer 2023-Present	<b>Yoonsoo Nam</b> , Masters, USC	
Summer 2023	<b>Ruyuan Zuo</b> , Masters, USC	<i>Now at Google</i>
Spring 2022	<b>Hanjie Chen / Junlin Wang / Abdallah Bashir</b> , PhD, UVA / Masters, UCI / Masters, Saarland Universit	<i>AI2 Intern</i>
Fall 2021	<b>Jillian Fisher / Liwei Jiang</b> , PhD, UW	
Summer 2021	<b>Kawin Ethayarajh</b> , PhD, Stanford University	<i>AI2 Intern</i>
Summer 2021	<b>Sarah Wiegrefe</b> , PhD, George Institute of Technology	<i>AI2 Intern</i>
Summer 2021	<b>Ximing Liu</b> , BS, UW CSE	<i>AI2 Intern</i>
Fall 2020-2023	<b>Alisa Liu</b> , PhD, UW CSE	
Fall 2020	<b>Alon Jacovi</b> , PhD, Bar Ilan University	<i>AI2 Intern</i>
Summer 2020-1	<b>Jenny Liang</b> , BS, UW CSE	<i>AI2 Intern</i>
Summer 2020	<b>Jize Cao</b> , BS, UW CSE	<i>AI2 Intern</i>
Spring 2020-2023	<b>Xuhui Zhou</b> , CLMS, UW Linguistics	
Fall 2019 - Spring 2020	<b>Yiben Yang</b> , PhD, Northwestern University	
Fall 2019 - Spring 2020	<b>Chaitanya Malaviya</b> , PYI, Allen Institute for AI	
2018 - 2019	<b>Ron Fan / Karishma Mandyam</b> , MS, UW CSE / BS, UW CSE	

## Teaching

### CLASSES

Spring 2024	<b>CSCI 499: Language Models in NLP</b> , Students: 28	<b>USC</b>
Fall 2023	<b>CSCI 499: Language Models in NLP</b> , Students: 30	<b>USC</b>
Fall 2022	<b>CSCI 699: Data-Centric NLP</b> , Students: 29	<b>USC</b>

### TUTORIALS

2019	<b>Transfer Learning in Natural Language Processing</b> , Attendees: 350+	<b>NAACL</b>
2018	<b>Frame Semantics across Languages: Towards a Multilingual FrameNet</b> , Attendees: 30+	<b>CoLING</b>

### GUEST LECTURES

Spring 2024	<b>Data-Centric NLP</b> , Host: Xuezhe Ma	<i>USC CSCI 544: Advanced NLP</i>
Summer 2023	<b>Introduction to NLP</b> , Host: Darin Gray	<i>Introduction to Engineering for USC</i> <i>Viterbi K-12 Discover Engineering</i>
Spring 2023	<b>Contextualizing Bias</b> , Host: Maja Matarić	<i>USC CSCI 697: Seminar in CS Research</i>
Spring 2022	<b>What's in a dataset?</b> , Host: Cyrus Shahabi	<i>USC CSCI 697: Seminar in CS Research</i>
Spring 2022	<b>Data-Centric NLP</b> , Host: Dongyeop Kang	<i>UMinn. CSCI 8980: Intro / NLP Research</i>
Spring 2022	<b>Annotating Ambiguous NLI</b> , Host: Dirk Hovy	<i>MilaNLP: Coding Aperitivo</i>
Spring 2021	<b>Transfer Learning</b> , Host: Zaid Harachoui	<i>UW DATA 598: Statistical Deep Learning</i>
Fall 2020	<b>Biases and Interpretability in NLP</b> , Host: Eunsol Choi	<i>UT-Austin CS 395T: Topics in NLP</i>
Winter 2019	<b>Phrase-Structure Parsing</b> , Host: Noah A. Smith	<i>UW CSE 447/547M: NLP</i>
Spring 2018	<b>Minimum Bayes Risk Decoding</b> , Host: Waleed Ammar	<i>UW CSE 599 D1: Advanced Topics in NLP</i>
Spring 2018	<b>Dependency Parsing</b> , Host: Noah A. Smith	<i>UW CSEP 517: NLP</i>
Fall 2017	<b>Unsupervised Learning</b> , Host: Noah A. Smith	<i>UW CSE 446: Machine Learning</i>
Fall 2014	<b>Dependency Parsing with Chu-Liu-Edmonds</b> , Hosts: Chris Dyer, Alon Lavie, Bob Frederking	<i>CMU 11-271: Algorithms for NLP</i>

### TEACHING ASSITANTSHIP

## Invited Talks

Feb 13, 2024	<b>LLM Inference for Scaling Data Creation</b> , Intel Rising Stars Talk	<b>Intel Labs</b>
Feb 8, 2024	<b>Understanding LLMs via their Generative Successes and Shortcomings</b> , LTL Seminar	<b>University of Cambridge</b>
Dec 15, 2023	<b>Understanding LLMs via their Generative Successes and Shortcomings</b> , Keynote	<b>ATTRIB Workshop @ NeurIPS 2023</b>
Nov 16, 2023	<b>Understanding Online Discourse through Social Context and Structured Pragmatics</b> , UCLA Communications Symposium	<b>UCLA</b>
Oct 3, 2023	<b>Understanding Data with <math>\mathcal{V}</math>-Information</b> , Keynote	<b>Towards the Next Generation of Computer Vision Datasets Workshop @ ICCV 2023</b>
Sep 22, 2023	<b>The Role of Language Models in Natural Language Processing</b> , Los Angeles County of Education Computer Science Speaker Series	<b>Los Angeles</b>
Aug 24, 2023	<b>Understanding Datasets and Explanations through <math>\mathcal{V}</math>-Information</b> , Alumni Keynote: CMU LTI Student Research Symposium	<b>CMU, Pittsburgh</b>
Jul 13, 2023	<b>Contextualizing Representations in Varied Annotator Perspectives</b> , Rep4NLP Workshop	<b>ACL Toronto</b>
Jun 14, 2023	<b>Contextualizing Data in Varied Annotator Perspectives</b> , Responsible Machine Learning	<b>Google Brain</b>
Mar 14, 2023	<b>Understanding Dataset Difficulty with <math>\mathcal{V}</math>-Usable Information</b> , Cohere for AI Reading Group	<b>Cohere for AI</b>
Feb 28, 2023	<b>Designing Controls and Filters for Dataset Generation</b> , Spotify Research Seminar	<b>Spotify Research Labs</b>
Feb 03, 2023	<b>What's in a Dataset? Interpreting datasets to enable better data creation</b> , Amazon Data-Centric AI Seminar	<b>Amazon</b>
Feb 01, 2023	<b>Contextualizing Bias in Hate Speech Detection</b> , CAIS++ Seminar	<b>USC</b>
Nov 18, 2022	<b>Generating Datasets for Robust Generalization</b> , SoCal NLP Symposium	<b>UC Santa Barbara</b>
Nov 09, 2022	<b>Contextualizing Bias in Hate Speech Detection through Annotator Perspectives</b> , Center for AI in Society (CAIS) Seminar	<b>USC</b>
May 23, 2022	<b>The Devil's in the Data: Mapping and Generating Datasets for Robust Generalization</b> , ACL Spotlight Talks for Young Rising Stars	<b>Dublin, Ireland</b>
May 13, 2022	<b>Mapping and Generating Datasets for Robust Generalization</b> , CS Seminar	<b>UC Irvine</b>
Mar 16, 2022	<b>Rethinking Dataset Construction: : The Role of Generative Modeling and Annotator Perspectives</b> , IFDS Ethics Seminar	<b>UC Santa Cruz</b>
Oct 20, 2021	<b>What's in your Data? Mapping Datasets and Exploring Data Usability</b> , Machine Learning Research Group Talks	<b>Oracle</b>
Feb 24, 2020	<b>Addressing Biases for Robust, Generalizable AI</b> , NERT Seminar	<b>Georgetown University</b>
Feb 12, 2020	<b>Addressing Biases for Robust, Generalizable AI</b> , NLP Seminar	<b>Georgia Tech</b>
Nov 02, 2020	<b>Responsible AI: Addressing Biases in Datasets and Models</b> , E+D Product Leads	<b>Microsoft</b>
Oct 30, 2019	<b>Biases and Learning Challenges in Natural Language Processing</b> , Rising Stars in EECS	<b>UIUC</b>
Oct 18, 2019	<b>Sprucing up a Dataset: Adversarially Filtering Dataset Artifacts</b> , Linguistics Colloquium	<b>UW Linguistics</b>
Apr 08, 2019	<b>Learning Challenges in Natural Language Processing</b> , AI2 Seminar	<b>Allen Institute for AI</b>
Apr 04, 2019	<b>Learning Challenges in Natural Language Processing</b> , NLP Seminar	<b>Stanford University</b>
Mar 13, 2019	<b>Learning Challenges in Natural Language Processing</b> , CLSP Seminar	<b>Johns Hopkins University</b>
Mar 11, 2019	<b>Learning Challenges in Natural Language Processing</b> , CILVR Seminar	<b>New York University</b>
Jan 22, 2019	<b>Learning Challenges in Natural Language Processing</b> , MSR Seminar	<b>Microsoft Research, New York</b>
Sep 15, 2018	<b>Representation Learning with Linguistic Structure</b> , MSR AI Breakthroughs Workshop	<b>Microsoft Research, Redmond</b>
Apr 27, 2018	<b>Syntactic Scaffolds for Semantic Structures</b> , NorthWest NLP	<b>Microsoft Research, Redmond</b>

## Invited Panels

Jul 13, 2023 **Limitations of Large Language Models**, Rep4NLP Workshop  
 Jun 29, 2023 **Generative AI and Education**, Roundtable  
 Apr 5, 2023 **Almaginings**, Polymathic Pizza  
 Jul 14, 2022 **Adversarial Data Augmentations**, DADC Workshop  
 Jun 17, 2022 **Role of LLMs**, Responsible AI Symposium

**ACL Toronto**  
**ASEE 2023**  
**USC Sidney Harmon Academy**  
**for Polymathic Studies**  
**NAACL Seattle**  
**AILA**

## Publications

### CONFERENCE PAPERS

- Closing the Curious Case of Neural Text Degeneration *ICLR (To Appear) 2024*  
 • M. Finlayson, J. Hewitt, A. Koller, **S. Swayamdipta**, A. Sabharwal
- Does Video Summarization Require Videos? Quantifying the Effectiveness of Language in Video Summarization *ICASSP (To Appear) 2023*  
 • Y. Nam, A. Lehari, D. Yang, D. Bose, **S. Swayamdipta**, and S. Narayanan
- We're Afraid Language Models Aren't Modeling Ambiguity *EMNLP 2023*  
 • A. Liu, Z. Wu, J. Michael, A. Suhr, P. West, A. Koller, **S. Swayamdipta**, N. A. Smith, and Y. Choi
- MAUVE Scores for Generative Models: Theory and Practice *JMLR 2023*  
 • K. Pillutla, L. Liu, J. Thickstun, S. Welleck, **S. Swayamdipta**, R. Zellers, S. Oh, Y. Choi, and Z. Harchaoui
- I2D2: Inductive Knowledge Distillation with NeuroLogic and Self-Imitation *ACL 2023*  
 • C. Bhagavatula, J. D. Hwang, D. Downey, R. Le Bras, X. Lu, K. Sakaguchi, **S. Swayamdipta**, P. West, and Y. Choi
- REV: Information-Theoretic Evaluation of Free-Text Rationales *ACL 2023*  
 • H. Chen, F. Brahman, X. Ren, Y. Ji, Y. Choi, and **S. Swayamdipta**
- COBRA Frames: Contextual Reasoning about Effects and Harms of Offensive Statements *ACL-Findings 2023*  
 • X. Zhou, H. Zhu, A. Yerukola, T. Davidson, J. D. Hwang, **S. Swayamdipta**, and M. Sap
- Investigating the Benefits of Free-Form Rationales *EMNLP-Findings 2022*  
 J. Sun, **S. Swayamdipta**, J. May, and X. Ma
- NeuroCounterfactuals: Beyond Minimal-Edit Counterfactuals for Richer Data Augmentation *EMNLP-Findings 2022*  
 P. Howard, G. Singer, V. Lal, Y. Choi, and **S. Swayamdipta**
- WaNLI: Worker and AI Collaboration for Natural Language Inference Dataset Creation *EMNLP-Findings 2022*  
 A. Liu, **S. Swayamdipta**, N. A. Smith, and Y. Choi
- Reframing Human-AI Collaboration for Generating Free-Text Explanations *NAACL 2022*  
 S. Wiegrefe, J. Hessel, **S. Swayamdipta**, M. Riedl, and Y. Choi
- Annotators with Attitudes: How Annotator Beliefs And Identities Bias Toxic Language Detection *NAACL 2022*  
 M. Sap, **S. Swayamdipta**, L. Vianna, X. Zhou, Y. Choi, and N. A. Smith
- Understanding Dataset Difficulty with V-Usable Information *ICML 2022*  
 • K. Ethayarajh, Y. Choi, and **S. Swayamdipta**  
 • **Outstanding Paper Award**
- MAUVE: Measuring the Gap Between Neural Text and Human Text using Divergence Frontiers *NeurIPS 2021*  
 • K. Pillutla, **S. Swayamdipta**, R. Zellers, J. Thickstun, S. Wellecks, Y. Choi, and Z. Harchaoui  
 • **Outstanding Paper Award**
- Contrastive Explanations for Model Interpretability *EMNLP 2021*  
 A. Jacovi, **S. Swayamdipta**, S. Ravfogel, Y. Elazar, Y. Choi, and Y. Goldberg
- On-the-Fly Controlled Text Generation with Experts and Anti-Experts *ACL 2021*  
 A. Liu, M. Sap, X. Lu, **S. Swayamdipta**, C. Bhagavatula, N. A. Smith, and Y. Choi
- Challenges in Automated Debiasing for Toxic Language Detection *EACL 2021*  
 • X. Zhou, M. Sap, **S. Swayamdipta**, N. A. Smith, and Y. Choi

Dataset Cartography: Mapping and Diagnosing Datasets with Training Dynamics	EMNLP 2020
• <b>S. Swayamdipta</b> , R. Schwartz, N. Lourie, Y. Wang, H. Hajishirzi, N. A. Smith, Y. Choi	
Generative Data Augmentation for Commonsense Reasoning	EMNLP-Findings 2020
• Y. Yang, C. Malaviya, J. Fernandez, <b>S. Swayamdipta</b> , R. LeBras, J. Wang, C. Bhagavatula, Y. Choi, and D. Downey	
Adversarial Filters of Dataset Biases	ICML 2020
• R. LeBras, <b>S. Swayamdipta</b> , C. Bhagavatula, R. Zellers, M. E. Peters, A. Sabharwal, and Y. Choi	
The Right Tool for the Job: Matching Model and Instance Complexities	ACL 2020
• R. Schwartz, G. Stanovsky, <b>S. Swayamdipta</b> , J. Dodge, and N. A. Smith	
Don't Stop Pretraining: Adapt Language Models to Domains and Tasks	ACL 2020
• S. Gururangan, A. Marasović, <b>S. Swayamdipta</b> , K. Lo, I. Beltagy, D. Downey, and N. A. Smith	
• <b>Best Paper Honorable Mention</b>	
Syntactic Scaffolds for Semantic Structures	EMNLP 2018
• <b>S. Swayamdipta</b> , S. Thomson, K. Lee, L. Zettlemoyer, C. Dyer, N. A. Smith.	
Learning Joint Semantic Parsers from Disjoint Data	NAACL 2018
• H. Peng, S. Thomson, <b>S. Swayamdipta</b> , N. A. Smith	
Annotation Artifacts in Natural Language Inference Data	NAACL 2018
• S. Gururangan*, <b>S. Swayamdipta</b> *, O. Levy, R. Schwartz, S. Bowman, and N. A. Smith	
• * equal contribution	
Polyglot Semantic Role Labeling	ACL 2018
• P. Mulcaire, <b>S. Swayamdipta</b> , and N. A. Smith	
Multi-Mention Learning for Reading Comprehension with Neural Cascades	ICLR 2018
• <b>S. Swayamdipta</b> , A. Parikh, T. Kwiatkowski	
Greedy, Joint Syntactic and Semantic Parsing with Stack LSTMs	CoNLL 2016
• <b>S. Swayamdipta</b> , M. Ballesteros, C. Dyer, N. A. Smith	
A Dependency Parser for Tweets	EMNLP 2014
• L. Kong, N. Schneider, <b>S. Swayamdipta</b> , A. Bhatia, C. Dyer, N. A. Smith	
The Pursuit of Power and its Manifestation in Written Dialog	ICSC 2012
• <b>S. Swayamdipta</b> , O. Rambow	

## WORKSHOP PAPERS

Sister Help: Data Augmentation for Frame-Semantic Role Labeling	LAW-DMR @ EMNLP 2021
• A. Pancholy, <b>S. Swayamdipta</b> , M. R. L. Petruck	
Multi-Task Learning for Incremental Parsing using Stack LSTMs	WiML @ NeurIPS 2016
• <b>S. Swayamdipta</b> , M. Ballesteros, C. Dyer, N. A. Smith	
CMU: Arc-Factored, Discriminative Semantic Dependency Parsing	SemEval 2014
• S. Thomson, D. Bamman, J. Dodge, <b>S. Swayamdipta</b> , N. Schneider, C. Dyer, N. A. Smith	
The CMU Machine Translation Systems	WMT 2014
• A. Matthews, C. Dyer, A. Lavie, G. Hanneman, W. Ammar, A. Bhatia, <b>S. Swayamdipta</b> , E. Schlinger, Y. Tsvetkov	

## WORKING PAPERS

Entropy-Constrained Non-Negative Kernel Regression for Sentence-Level Anomaly Detection	Under Submission
• A. Gulati, X. Dong, C. Hurtado, S. Shekkizhar, <b>S. Swayamdipta</b> , A. Ortega	
OATH-Frames: Characterizing Online Attitudes towards Homelessness via LLM Assistants	Under Submission
• J. Ranjit, B. Joshi, R. Dorn, L. Petry, O. Koumoundouros, J. Bottarini, P. Liu, E. Rice, and <b>S. Swayamdipta</b>	
Annotating FrameNet via Structure-Conditioned Language Generation	Under Submission
• X. Cui, and <b>S. Swayamdipta</b>	
NeuroComparatives: Neuro-Symbolic Distillation of Comparative Knowledge	arXiv:2305.04978
• P. Howard, J. Wang, V. Lal, G. Singer, Y. Choi, and <b>S. Swayamdipta</b>	

Shallow Syntax in Deep Water

[arXiv:1908.11047](#)

- **S. Swayamdipta**, M. Peters, B. Roof, C. Dyer and N. A. Smith

Frame-Semantic Parsing with Softmax-Margin Segmental RNNs and a Syntactic Scaffold

[arXiv:1706.09528](#)

- **S. Swayamdipta**, S. Thomson, C. Dyer and N. A. Smith

DyNet: The Dynamic Neural Network Toolkit

[arXiv:1701.03980](#)

- G. Neubig, C. Dyer, Y. Goldberg, A. Matthews, W. Ammar, A. Anastasopoulos, M. Ballesteros, D. Chiang, D. Clothiaux, T. Cohn, K. Duh, M. Faruqui, C. Gan, D. Garrette, Y. Ji, L. Kong, A. Kuncoro, G. Kumar, C. Malaviya, P. Michel, Y. Oda, M. Richardson, N. Saphra, **S. Swayamdipta**, P. Yin

## TUTORIAL AND ORGANIZATION PAPERS

Proceedings of the 3rd Workshop on Deep Learning Approaches for Low-Resource NLP (DeepLo 2022)

[NAACL 2022](#)

- C. Cherry, A. Fan, G. Foster, G. Haffari, S. Khadivi, N. Peng, X. Ren, E. Shareghi, **S. Swayamdipta**

Transfer Learning in Natural Language Processing

[NAACL 2019](#)

- S. Ruder, M. E. Peters, **S. Swayamdipta**, T. Wolf

Proceedings of the 2nd Workshop on Deep Learning Approaches for Low-Resource NLP (DeepLo 2019)

[EMNLP 2019](#)

- C. Cherry, G. Durrett, G. Foster, R. Haffari, S. Khadivi, N. Peng, X. Ren, **S. Swayamdipta**

Frame Semantics across Languages: Towards a Multilingual FrameNet

[CoLING 2018](#)

- C. F. Baker, M. Ellsworth, M. R. L. Petruck, **S. Swayamdipta**

## Professional Service

---

### EXTERNAL SERVICE

#### Senior Area Chair

- ACL 2024: Sentence-level Semantics

#### Area Chair

- COLM 2024 (Upcoming)
- ICLR 2024
- EMNLP 2023: Machine Learning for NLP
- ACL 2023: Interpretability
- EMNLP 2022: Language Models
- EMNLP 2021: Machine Learning for NLP
- NAACL 2021: Sentence-Level Semantics
- EACL 2021: Sentence-Level Semantics
- ACL 2020: Semantics (Long)

#### Reviewer: Conferences

- |             |                        |
|-------------|------------------------|
| • ACL (ARR) | 2015-2022              |
| • NAACL     | 2015-2022              |
| • EMNLP     | 2015-2022              |
| • NeurIPS   | 2018-2022              |
| • ICML      | 2015, 2019, 2020, 2023 |
| • EACL      | 2017                   |
| • AAAI      | 2017-2020              |
| • CoNLL     | 2017-2018, 2020        |

#### Reviewer: Workshops

- |  |      |
|--|------|
| • NAACL Student Research Workshop                          | 2016 |
| • Workshop for Women in Machine Learning (WiML) at NeurIPS | 2016 |
| • Machine Reading for Question Answering (MRQA) Workshop   | 2017 |

#### Reviewer: Journals

- |                             |                |
|-----------------------------|----------------|
| • Transactions of ACL       | 2020 - PRESENT |
| • Computational Linguistics | 2019 - 2022    |
| • Journal of AI Research    | 2019           |

## Co-organizer

- EACL 2024: Workshop on Uncertainty in NLP 2024
- NAACL 2022: Workshop on Deep Learning in Low Resource NLP (DeepLo) 2022
- EMNLP 2019: Workshop on Deep Learning in Low Resource NLP (DeepLo) 2019
- West Coast NLP Workshop 2018
- Women's Research Day 2017
- UW-NLP Retreat 2015-2016

## Mentorship at Conferences

- EMNLP 2020, ACL 2020, EMNLP 2018

## INTERNAL SERVICE

### Univeristy of Southern California

*Los Angeles, CA, USA*

- Diversity, Equity and Inclusion Committee 2023-present
- Distinguished Lectures Committee 2023-present
- Annual Faculty Review Committee 2023-present
- PhD Fellowship Admissions Committee 2022-2023
- Viterbi Undergraduate Admission Scholarship 2023-2024

### Allen Institute for AI

*Seattle WA, USA*

- Diversity, Equity and Inclusion Committee 2019-2022
- AllenAI Outstanding Engineer Scholarship Committee 2019

### Paul G. Allen School of CSE, University of Washington

*Seattle WA, USA*

- Faculty Recruiting Liaison 2019
- Graduate Student Advisory Council 2018
- Outreach: Highline Public High School 2017
- CS: Gradswomen 2016-2019

## Internships

### Allen Institute for AI

*Seattle, WA, USA*

RESEARCH INTERN, **Host:** MATTHEW E. PETERS

*Summer-Fall 2018*

- **Project:** Shallow Syntactic Priming of Large-scale Language Models

### Google Inc.

*New York, NY, USA*

RESEARCH INTERN, **Host:** ANKUR P. PARIKH

*Summer 2017*

- **Project:** Machine Comprehension for Google Neon

### Columbia University

*New York, NY, USA*

RESEARCH ASSISTANT, CENTER FOR COMPUTATIONAL LEARNING SYSTEMS

*Jan 2013 - May 2013*

### Google Inc.

*Mountain View, CA, USA*

SOFTWARE ENGINEER INTERN

*Summer 2012*

- Information retrieval evaluation for Google TV search

### Microsoft Corporation

*Hyderabad, India*

SOFTWARE DEVELOPMENT INTERN

*Summer 2009*

- **Project:** Search functionality for customer hierarchies

## References

**Yejin Choi**, yejin@cs.washington.edu

*UW; AI2*

**Noah A. Smith**, noah@cs.washington.edu

*UW; AI2*

**Chris Dyer**, cdyer@google.com

*Google DeepMind*

**Luke Zettlemoyer**, lsz@cs.washington.edu

*UW; FAIR*