

# Swabha Swayamdipta

GABILAN ASSISTANT PROFESSOR OF COMPUTER SCIENCE, UNIVERSITY OF SOUTHERN CALIFORNIA

✉ swabhas@usc.edu | 🏠 https://swabhs.com | 🌐 swabhs | 🐦 @swabh

## Professional Experience

### University of Southern California

GABILAN ASSISTANT PROFESSOR OF COMPUTER SCIENCE

- Associate Director of the Center for AI and Society, since Fall 2023

Los Angeles, CA, USA

Aug 2022 - PRESENT

### Allen Institute for AI

POSTDOCTORAL INVESTIGATOR

- **Advisor:** Yejin Choi
- **Project Team:** MOSAIC

Seattle, WA, USA

Aug 2019 - June 2022

### Oracle Server Technologies

MEMBER TECHNICAL STAFF

- **Project:** Unified messaging for Fusion Middleware

Bangalore, India

2010 - 2011

## Education

### Carnegie Mellon University

PHD IN LANGUAGE AND INFORMATION TECHNOLOGIES

- **Advisors:** Noah A. Smith, Chris Dyer
- **University of Washington**, Seattle, WA — *Visiting PhD student Fall 2015 - Summer 2019*
- **Thesis:** Syntactic Inductive Biases in NLP

Pittsburgh, PA, USA

2013 - 2019

### Columbia University

MASTERS IN COMPUTER SCIENCE

- **Advisors:** Owen Rambow, Michael Collins
- **Specialization:** Natural Language Processing

City of New York, NY, USA

2011 - 2012

### National Institute of Technology

BACHELORS OF TECHNOLOGY IN COMPUTER SCIENCE AND ENGINEERING

Calicut, Kerala, India

2006 - 2010

## Awards and Grants

2024	<b>Simons Institute Research Fellowship</b> , Special Program on LLMs and Transformers, Spring 2025	UC Berkeley
2024	<b>NSF CISE Core Medium</b> , RII: Hard Data to the Model	NSF
2024	<b>Zumberge Fellowship</b> , DEI in Research	USC OORI
2023	<b>Intel Rising Stars Award</b> , LLM Inference for Scaling Data Creation	Intel Labs
2022	<b>AI2 Young Investigators Award</b> , Synthesizing Data Towards Commonsense and Causal Reasoning	Allen Institute for AI
2022	<b>Gabilan Assistant Professor Fellowship</b> ,	USC WISE
2022	<b>Outstanding Paper Award</b> , Understanding Dataset Difficulty with V-Usable Information	ICML
2021	<b>Outstanding Paper Award</b> , MAUVE: Measuring the Gap Between Neural Text and Human Text using Divergence Frontiers	NeurIPS
2020	<b>Best Paper: Honorable Mention</b> , Don't Stop Pretraining: Adapt Language Models to Domains and Tasks	ACL
2019	<b>Rising Star in EECS</b> , UIUC	
2014	<b>Best Student Presentation</b> , Diversity in Dependency Parsing	CMU LTI Student Research Symposium
2009	<b>Sun Campus Ambassador Scholarship</b> , Bangalore, India	Sun Microsystems
2003	<b>National Talent Search Scholarship</b> , Rourkela, India	NCERT, India

## Advising

Fall 2024-Present	<b>Gregory Yauney</b> , PostDoc, USC	
Fall 2024-Present	<b>Xinyue Cui</b> , PhD, USC	
Fall 2024-Present	<b>Muru Zhang</b> , PhD, USC	<i>Co-Advisor: Robin Jia</i>
Fall 2024-Present	<b>Atharva Kulkarni</b> , PhD, USC	
Spring 2024-Present	<b>Brihi Joshi</b> , PhD, USC	<i>Co-Advisor: Xiang Ren</i>
Fall 2023-Present	<b>Matthew Finlayson</b> , PhD, USC	<i>Co-Advisor: Xiang Ren</i>
Fall 2022-Present	<b>Sayan Ghosh</b> , PhD, USC	
Fall 2023-Present	<b>Jaspreet Ranjit</b> , PhD, USC	
Spring 2024-Present	<b>Shahzaib Saqib Warraich</b> , Masters, USC	
Spring 2024-Present	<b>Shauryasikt Jena</b> , Masters, USC	
Spring 2024-Present	<b>Xingjian Dong</b> , Masters, USC	
Spring 2024-Present	<b>Ryan Wang</b> , BS, USC	<i>Co-Advisor: Robin Jia</i>
Spring 2024-Present	<b>Joseph Liu</b> , BS, USC	
Spring 2024-Present	<b>Keyu He</b> , BS, USC	
Fall 2023-Present	<b>Risha Surana</b> , BS, USC	
Summer 2023-Present	<b>Aryan Gulati</b> , BS, USC	<i>Co-Advisor: Antonio Ortega</i>
Spring 2023-2024	<b>Xinyue Cui</b> , Masters, USC	<i>Now PhD@USC</i>
Fall 2023-Spring 2024	<b>Jacky Mo</b> , BS, USC	
Summer-Fall 2023	<b>Jayne Bottarini</b> , BS, USC	
Summer-Fall 2023	<b>Kritin Dhoka</b> , BS, USC	
Summer-Fall 2023	<b>Catherine He</b> , BS, USC	
Spring 2023-2024	<b>Yoonsoo Nam</b> , Masters, USC	<i>Now at RenderWolf</i>
Spring-Summer 2023	<b>Adam Lehavi</b> , BS, USC	
Summer 2023	<b>Ruyuan Zuo</b> , Masters, USC	<i>Now at Google</i>
Spring 2022	<b>Hanjie Chen / Junlin Wang / Abdallah Bashir</b> , PhD, UVA / Masters, UCI / Masters, Saarland Universit	<i>AI2 Intern</i>
Fall 2021	<b>Jillian Fisher / Liwei Jiang</b> , PhD, UW	
Summer 2021	<b>Kawin Ethayarajh</b> , PhD, Stanford University	<i>AI2 Intern</i>
Summer 2021	<b>Sarah Wiegrefe</b> , PhD, George Institute of Technology	<i>AI2 Intern</i>
Summer 2021	<b>Ximing Liu</b> , BS, UW CSE	<i>AI2 Intern</i>
Fall 2020-2023	<b>Alisa Liu</b> , PhD, UW CSE	
Fall 2020	<b>Alon Jacovi</b> , PhD, Bar Ilan University	<i>AI2 Intern</i>
Summer 2020-1	<b>Jenny Liang</b> , BS, UW CSE	<i>AI2 Intern</i>
Summer 2020	<b>Jize Cao</b> , BS, UW CSE	<i>AI2 Intern</i>
Spring 2020-2023	<b>Xuhui Zhou</b> , CLMS, UW Linguistics	
Fall 2019 - Spring 2020	<b>Yiben Yang</b> , PhD, Northwestern University	
Fall 2019 - Spring 2020	<b>Chaitanya Malaviya</b> , PYI, Allen Institute for AI	
2018 - 2019	<b>Ron Fan / Karishma Mandyam</b> , MS, UW CSE / BS, UW CSE	

## Publications

### CONFERENCE PAPERS

Logits of API-Protected LLMs Leak Proprietary Information	<i>COLM 2024 (To Appear)</i>
• M. Finlayson, X. Ren, <b>S. Swayamdipta</b> ,	
Crowd-Calibrator: Can Annotator Disagreement Inform Calibration in Subjective Tasks?	<i>COLM 2024 (To Appear)</i>
• U. Khurana, E. Nalisnick, A. Fokkens, <b>S. Swayamdipta</b> ,	
Annotating FrameNet via Structure-Conditioned Language Generation	<i>ACL 2024 (To Appear)</i>
• X. Cui, and <b>S. Swayamdipta</b>	
Closing the Curious Case of Neural Text Degeneration	<i>ICLR 2024</i>
• M. Finlayson, J. Hewitt, A. Koller, <b>S. Swayamdipta</b> , A. Sabharwal	

Does Video Summarization Require Videos? Quantifying the Effectiveness of Language in Video Summarization	ICASSP 2024
• Y. Nam, A. Lehavi, D. Yang, D. Bose, <b>S. Swayamdipta</b> , and S. Narayanan	
NeuroComparatives: Neuro-Symbolic Distillation of Comparative Knowledge	NAACL-Findings 2024
• P. Howard, J. Wang, V. Lal, G. Singer, Y. Choi, and <b>S. Swayamdipta</b>	
Generative Explanations for Program Synthesizers	ICSE + VL/HCC 2024
• A. Nazari, S. Chattopadhyay, <b>S. Swayamdipta</b> , M. Raghothaman	
We're Afraid Language Models Aren't Modeling Ambiguity	EMNLP 2023
• A. Liu, Z. Wu, J. Michael, A. Suhr, P. West, A. Koller, <b>S. Swayamdipta</b> , N. A. Smith, and Y. Choi	
MAUVE Scores for Generative Models: Theory and Practice	JMLR 2023
• K. Pillutla, L. Liu, J. Thickstun, S. Welleck, <b>S. Swayamdipta</b> , R. Zellers, S. Oh, Y. Choi, and Z. Harchaoui	
I2D2: Inductive Knowledge Distillation with NeuroLogic and Self-Imitation	ACL 2023
• C. Bhagavatula, J. D. Hwang, D. Downey, R. Le Bras, X. Lu, K. Sakaguchi, <b>S. Swayamdipta</b> , P. West, and Y. Choi	
REV: Information-Theoretic Evaluation of Free-Text Rationales	ACL 2023
• H. Chen, F. Brahman, X. Ren, Y. Ji, Y. Choi, and <b>S. Swayamdipta</b>	
COBRA Frames: Contextual Reasoning about Effects and Harms of Offensive Statements	ACL-Findings 2023
• X. Zhou, H. Zhu, A. Yerukola, T. Davidson, J. D. Hwang, <b>S. Swayamdipta</b> , and M. Sap	
Investigating the Benefits of Free-Form Rationales	EMNLP-Findings 2022
J. Sun, <b>S. Swayamdipta</b> , J. May, and X. Ma	
NeuroCounterfactuals: Beyond Minimal-Edit Counterfactuals for Richer Data Augmentation	EMNLP-Findings 2022
P. Howard, G. Singer, V. Lal, Y. Choi, and <b>S. Swayamdipta</b>	
WaNLI: Worker and AI Collaboration for Natural Language Inference Dataset Creation	EMNLP-Findings 2022
A. Liu, <b>S. Swayamdipta</b> , N. A. Smith, and Y. Choi	
Reframing Human-AI Collaboration for Generating Free-Text Explanations	NAACL 2022
S. Wiegrefe, J. Hessel, <b>S. Swayamdipta</b> , M. Riedl, and Y. Choi	
Annotators with Attitudes: How Annotator Beliefs And Identities Bias Toxic Language Detection	NAACL 2022
M. Sap, <b>S. Swayamdipta</b> , L. Vianna, X. Zhou, Y. Choi, and N. A. Smith	
Understanding Dataset Difficulty with V-Usable Information	ICML 2022
• K. Ethayarajh, Y. Choi, and <b>S. Swayamdipta</b>	
• <b>Outstanding Paper Award</b>	
MAUVE: Measuring the Gap Between Neural Text and Human Text using Divergence Frontiers	NeurIPS 2021
• K. Pillutla, <b>S. Swayamdipta</b> , R. Zellers, J. Thickstun, S. Wellecks, Y. Choi, and Z. Harchaoui	
• <b>Outstanding Paper Award</b>	
Contrastive Explanations for Model Interpretability	EMNLP 2021
A. Jacovi, <b>S. Swayamdipta</b> , S. Ravfogel, Y. Elazar, Y. Choi, and Y. Goldberg	
On-the-Fly Controlled Text Generation with Experts and Anti-Experts	ACL 2021
A. Liu, M. Sap, X. Lu, <b>S. Swayamdipta</b> , C. Bhagavatula, N. A. Smith, and Y. Choi	
Challenges in Automated Debiasing for Toxic Language Detection	EACL 2021
• X. Zhou, M. Sap, <b>S. Swayamdipta</b> , N. A. Smith, and Y. Choi	
Dataset Cartography: Mapping and Diagnosing Datasets with Training Dynamics	EMNLP 2020
• <b>S. Swayamdipta</b> , R. Schwartz, N. Lourie, Y. Wang, H. Hajishirzi, N. A. Smith, Y. Choi	
Generative Data Augmentation for Commonsense Reasoning	EMNLP-Findings 2020
• Y. Yang, C. Malaviya, J. Fernandez, <b>S. Swayamdipta</b> , R. LeBras, J. Wang, C. Bhagavatula, Y. Choi, and D. Downey	
Adversarial Filters of Dataset Biases	ICML 2020
• R. LeBras, <b>S. Swayamdipta</b> , C. Bhagavatula, R. Zellers, M. E. Peters, A. Sabharwal, and Y. Choi	
The Right Tool for the Job: Matching Model and Instance Complexities	ACL 2020
• R. Schwartz, G. Stanovsky, <b>S. Swayamdipta</b> , J. Dodge, and N. A. Smith	

Don't Stop Pretraining: Adapt Language Models to Domains and Tasks	ACL 2020
<ul style="list-style-type: none"> <li>S. Gururangan, A. Marasović, <b>S. Swayamdipta</b>, K. Lo, I. Beltagy, D. Downey, and N. A. Smith</li> <li><b>Best Paper Honorable Mention</b></li> </ul>	
Syntactic Scaffolds for Semantic Structures	EMNLP 2018
<ul style="list-style-type: none"> <li><b>S. Swayamdipta</b>, S. Thomson, K. Lee, L. Zettlemoyer, C. Dyer, N. A. Smith.</li> </ul>	
Learning Joint Semantic Parsers from Disjoint Data	NAACL 2018
<ul style="list-style-type: none"> <li>H. Peng, S. Thomson, <b>S. Swayamdipta</b>, N. A. Smith</li> </ul>	
Annotation Artifacts in Natural Language Inference Data	NAACL 2018
<ul style="list-style-type: none"> <li>S. Gururangan*, <b>S. Swayamdipta</b>*, O. Levy, R. Schwartz, S. Bowman, and N. A. Smith</li> <li>* equal contribution</li> </ul>	
Polyglot Semantic Role Labeling	ACL 2018
<ul style="list-style-type: none"> <li>P. Mulcaire, <b>S. Swayamdipta</b>, and N. A. Smith</li> </ul>	
Multi-Mention Learning for Reading Comprehension with Neural Cascades	ICLR 2018
<ul style="list-style-type: none"> <li><b>S. Swayamdipta</b>, A. Parikh, T. Kwiatkowski</li> </ul>	
Greedy, Joint Syntactic and Semantic Parsing with Stack LSTMs	CoNLL 2016
<ul style="list-style-type: none"> <li><b>S. Swayamdipta</b>, M. Ballesteros, C. Dyer, N. A. Smith</li> </ul>	
A Dependency Parser for Tweets	EMNLP 2014
<ul style="list-style-type: none"> <li>L. Kong, N. Schneider, <b>S. Swayamdipta</b>, A. Bhatia, C. Dyer, N. A. Smith</li> </ul>	
The Pursuit of Power and its Manifestation in Written Dialog	ICSC 2012
<ul style="list-style-type: none"> <li><b>S. Swayamdipta</b>, O. Rambow</li> </ul>	

## WORKING PAPERS

Compare without Despair: Reliable Preference Evaluation with Generation Separability	Under Submission
<ul style="list-style-type: none"> <li>S. Ghosh, T. Srinivasan, <b>S. Swayamdipta</b>,</li> </ul>	
Out-of-Distribution Detection through Soft Clustering with Non-Negative Kernel Regression	Under Submission
<ul style="list-style-type: none"> <li>A. Gulati, X. Dong, C. Hurtado, S. Shekizhar, <b>S. Swayamdipta</b>, A. Ortega</li> </ul>	
OATH-Frames: Characterizing Online Attitudes towards Homelessness via LLM Assistants	Under Submission
<ul style="list-style-type: none"> <li>J. Ranjit, B. Joshi, R. Dorn, L. Petry, O. Koumoundouros, J. Bottarini, P. Liu, E. Rice, and <b>S. Swayamdipta</b></li> <li><b>Best Poster @ ShowCAIS 2024</b></li> </ul>	
Shallow Syntax in Deep Water	arXiv:1908.11047
<ul style="list-style-type: none"> <li><b>S. Swayamdipta</b>, M. Peters, B. Roof, C. Dyer and N. A. Smith</li> </ul>	
Frame-Semantic Parsing with Softmax-Margin Segmental RNNs and a Syntactic Scaffold	arXiv:1706.09528
<ul style="list-style-type: none"> <li><b>S. Swayamdipta</b>, S. Thomson, C. Dyer and N. A. Smith</li> </ul>	
DyNet: The Dynamic Neural Network Toolkit	arXiv:1701.03980
<ul style="list-style-type: none"> <li>G. Neubig, C. Dyer, Y. Goldberg, A. Matthews, W. Ammar, A. Anastasopoulos, M. Ballesteros, D. Chiang, D. Clothiaux, T. Cohn, K. Duh, M. Faruqi, C. Gan, D. Garrette, Y. Ji, L. Kong, A. Kuncoro, G. Kumar, C. Malaviya, P. Michel, Y. Oda, M. Richardson, N. Saphra, <b>S. Swayamdipta</b>, P. Yin</li> </ul>	

## WORKSHOP PAPERS

Sister Help: Data Augmentation for Frame-Semantic Role Labeling	LAW-DMR @ EMNLP 2021
<ul style="list-style-type: none"> <li>A. Pancholy, <b>S. Swayamdipta</b>, M. R. L. Petruck</li> </ul>	
Multi-Task Learning for Incremental Parsing using Stack LSTMs	WiML @ NeurIPS 2016
<ul style="list-style-type: none"> <li><b>S. Swayamdipta</b>, M. Ballesteros, C. Dyer, N. A. Smith</li> </ul>	
CMU: Arc-Factored, Discriminative Semantic Dependency Parsing	SemEval 2014
<ul style="list-style-type: none"> <li>S. Thomson, D. Bamman, J. Dodge, <b>S. Swayamdipta</b>, N. Schneider, C. Dyer, N. A. Smith</li> </ul>	
The CMU Machine Translation Systems	WMT 2014
<ul style="list-style-type: none"> <li>A. Matthews, C. Dyer, A. Lavie, G. Hanneman, W. Ammar, A. Bhatia, <b>S. Swayamdipta</b>, E. Schlinger, Y. Tsvetkov</li> </ul>	

## Invited Talks

---

Dec 5, 2024	<b>(Upcoming) Learning Machine Seminar Series (LMSS),</b>	<b>Cornell Tech</b>
Oct 10, 2024	<b>(Upcoming) NLP/Text-as-Data Talk Series,</b>	<b>New York University</b>
May 3, 2024	<b>Understanding LLMs through their Generative Behavior, Successes and Shortcomings,</b> ISI AI Seminar	<b>USC ISI</b>
Mar 26, 2024	<b>Towards (Closed-Source) LLM Accountability via Logit Signatures,</b> NSF-Open Source Generative AI Workshop	<b>Cornell Tech</b>
Mar 13, 2024	<b>Understanding LLMs through their Generative Behavior, Successes and Shortcomings,</b> Data Science Seminar	<b>University of Utah</b>
Feb 13, 2024	<b>LLM Inference for Scaling Data Creation,</b> Intel Rising Stars Talk	<b>Intel Labs</b>
Feb 8, 2024	<b>Understanding LLMs via their Generative Successes and Shortcomings,</b> LTL Seminar	<b>University of Cambridge</b>
Dec 15, 2023	<b>Understanding LLMs via their Generative Successes and Shortcomings,</b> Keynote	<b>ATTRIB Workshop @ NeurIPS 2023</b>
Nov 16, 2023	<b>Understanding Online Discourse through Social Context and Structured Pragmatics,</b> UCLA Communications Symposium	<b>UCLA</b>
Oct 3, 2023	<b>Understanding Data with <math>\mathcal{V}</math>-Information,</b> Keynote	<b>Towards the Next Generation of Computer Vision Datasets Workshop @ ICCV 2023</b>
Sep 22, 2023	<b>The Role of Language Models in Natural Language Processing,</b> Los Angeles County of Education Computer Science Speaker Series	<b>Los Angeles</b>
Aug 24, 2023	<b>Understanding Datasets and Explanations through <math>\mathcal{V}</math>-Information,</b> Alumni Keynote: CMU LTI Student Research Symposium	<b>CMU, Pittsburgh</b>
Jul 13, 2023	<b>Contextualizing Representations in Varied Annotator Perspectives,</b> Rep4NLP Workshop	<b>ACL Toronto</b>
Jun 14, 2023	<b>Contextualizing Data in Varied Annotator Perspectives,</b> Responsible Machine Learning	<b>Google Brain</b>
Mar 14, 2023	<b>Understanding Dataset Difficulty with <math>\mathcal{V}</math>-Usable Information,</b> Cohere for AI Reading Group	<b>Cohere for AI</b>
Feb 28, 2023	<b>Designing Controls and Filters for Dataset Generation,</b> Spotify Research Seminar	<b>Spotify Research Labs</b>
Feb 03, 2023	<b>What's in a Dataset? Interpreting datasets to enable better data creation,</b> Amazon Data-Centric AI Seminar	<b>Amazon</b>
Feb 01, 2023	<b>Contextualizing Bias in Hate Speech Detection,</b> CAIS++ Seminar	<b>USC</b>
Nov 18, 2022	<b>Generating Datasets for Robust Generalization,</b> SoCal NLP Symposium	<b>UC Santa Barbara</b>
Nov 09, 2022	<b>Contextualizing Bias in Hate Speech Detection through Annotator Perspectives,</b> Center for AI in Society (CAIS) Seminar	<b>USC</b>
May 23, 2022	<b>The Devil's in the Data: Mapping and Generating Datasets for Robust Generalization,</b> ACL Spotlight Talks for Young Rising Stars	<b>Dublin, Ireland</b>
May 13, 2022	<b>Mapping and Generating Datasets for Robust Generalization,</b> CS Seminar	<b>UC Irvine</b>
Mar 16, 2022	<b>Rethinking Dataset Construction: : The Role of Generative Modeling and Annotator Perspectives,</b> IFDS Ethics Seminar	<b>UC Santa Cruz</b>
Oct 20, 2021	<b>What's in your Data? Mapping Datasets and Exploring Data Usability,</b> Machine Learning Research Group Talks	<b>Oracle</b>
Feb 24, 2020	<b>Addressing Biases for Robust, Generalizable AI,</b> NERT Seminar	<b>Georgetown University</b>
Feb 12, 2020	<b>Addressing Biases for Robust, Generalizable AI,</b> NLP Seminar	<b>Georgia Tech</b>
Nov 02, 2020	<b>Responsible AI: Addressing Biases in Datasets and Models,</b> E+D Product Leads	<b>Microsoft</b>
Oct 30, 2019	<b>Biases and Learning Challenges in Natural Language Processing,</b> Rising Stars in EECS	<b>UIUC</b>
Oct 18, 2019	<b>Sprucing up a Dataset: Adversarially Filtering Dataset Artifacts,</b> Linguistics Colloquium	<b>UW Linguistics</b>
Apr 08, 2019	<b>Learning Challenges in Natural Language Processing,</b> AI2 Seminar	<b>Allen Institute for AI</b>
Apr 04, 2019	<b>Learning Challenges in Natural Language Processing,</b> NLP Seminar	<b>Stanford University</b>
Mar 13, 2019	<b>Learning Challenges in Natural Language Processing,</b> CLSP Seminar	<b>Johns Hopkins University</b>
Mar 11, 2019	<b>Learning Challenges in Natural Language Processing,</b> CILVR Seminar	<b>New York University</b>
Jan 22, 2019	<b>Learning Challenges in Natural Language Processing,</b> MSR Seminar	<b>Microsoft Research, New York</b>
Sep 15, 2018	<b>Representation Learning with Linguistic Structure,</b> MSR AI Breakthroughs Workshop	<b>Microsoft Research, Redmond</b>
Apr 27, 2018	<b>Syntactic Scaffolds for Semantic Structures,</b> NorthWest NLP	<b>Microsoft Research, Redmond</b>

## Invited Panels

Oct 18, 2024 **(Upcoming) Alfred P. Sloan Foundation Science Seminar,**  
 Jul 13, 2023 **Limitations of Large Language Models,** Rep4NLP Workshop  
 Jun 29, 2023 **Generative AI and Education,** Roundtable  
 Apr 5, 2023 **Almaginings,** Polymathic Pizza  
 Jul 14, 2022 **Adversarial Data Augmentations,** DADC Workshop  
 Jun 17, 2022 **Role of LLMs,** Responsible AI Symposium

*USC School of Cinematic Arts*  
*ACL Toronto*  
*ASEE 2023*  
*USC Sidney Harmon Academy*  
*for Polymathic Studies*  
*NAACL Seattle*  
*AILA*

## Professional Service

### EXTERNAL SERVICE

#### Senior Area Chair

- EMNLP 2024: Machine Learning in NLP
- ACL 2024: Sentence-level Semantics

#### Area Chair

- ICLR 2025 (upcoming)
- COLM 2024
- ICLR 2024
- EMNLP 2023: Machine Learning for NLP
- ACL 2023: Interpretability
- EMNLP 2022: Language Models
- EMNLP 2021: Machine Learning for NLP
- NAACL 2021: Sentence-Level Semantics
- EACL 2021: Sentence-Level Semantics
- ACL 2020: Semantics (Long)

#### Reviewer: Conferences

- ACL (ARR) 2015-2022
- NAACL 2015-2022
- EMNLP 2015-2022
- NeurIPS 2018-2022
- ICML 2015, 2019, 2020, 2023
- EACL 2017
- AAAI 2017-2020
- CoNLL 2017-2018, 2020

#### Reviewer: Workshops

- NAACL Student Research Workshop 2016
- Workshop for Women in Machine Learning (WiML) at NeurIPS 2016
- Machine Reading for Question Answering (MRQA) Workshop 2017

#### Reviewer: Journals

- Transactions of ACL 2020 - PRESENT
- Computational Linguistics 2019 - 2022
- Journal of AI Research 2019

#### Co-organizer

- EACL 2024: Workshop on Uncertainty in NLP 2024
- NAACL 2022: Workshop on Deep Learning in Low Resource NLP (DeepLo) 2022
- EMNLP 2019: Workshop on Deep Learning in Low Resource NLP (DeepLo) 2019
- West Coast NLP Workshop 2018
- Women's Research Day 2017
- UW-NLP Retreat 2015-2016

#### Mentorship at Conferences

- EMNLP 2020, ACL 2020, EMNLP 2018

### INTERNAL SERVICE



## Univeristy of Southern California

- Diversity, Equity and Inclusion Committee
- Apple and IBM Fellowship Committee
- Distinguished Lectures Committee
- Annual Faculty Review Committee
- Apple and IBM Fellowship Committee
- PhD Fellowship Admissions Committee
- Viterbi Undergraduate Admission Scholarship

Los Angeles, CA, USA

2023-present  
2024  
2023-2024  
2023-2024  
2023  
2022-2023  
2023-2024

## Allen Institute for AI

- Diversity, Equity and Inclusion Committee
- AllenAI Outstanding Engineer Scholarship Committee

Seattle WA, USA

2019-2022  
2019

## Paul G. Allen School of CSE, University of Washington

- Faculty Recruiting Liaison
- Graduate Student Advisory Council
- Outreach: Highline Public High School
- CS: Gradswomen

Seattle WA, USA

2019  
2018  
2017  
2016-2019

## Teaching

### CLASSES

Fall 2024	<b>CSCI 544: Applied NLP</b> , Students: 260	<b>USC</b>
Spring 2024	<b>CSCI 499: Language Models in NLP</b> , Students: 28	<b>USC</b>
Fall 2023	<b>CSCI 499: Language Models in NLP</b> , Students: 30	<b>USC</b>
Fall 2022	<b>CSCI 699: Data-Centric NLP</b> , Students: 29	<b>USC</b>

### TUTORIALS

2022	<b>OAMLS: Online Asian Machine Learning School</b> , Attendees: 350+	<b>ACML</b>
2019	<b>Transfer Learning in Natural Language Processing</b> , Attendees: 350+	<b>NAACL</b>
2018	<b>Frame Semantics across Languages: Towards a Multilingual FrameNet</b> , Attendees: 30+	<b>CoLING</b>

### TEACHING ASSISTANTSHIP

Spring 2015	<b>CMU 11-411/611: Natural Language Processing</b> , Instructors: Chris Dyer, Alan W Black, Shomir Wilson	<b>Head TA</b>
-------------	---	----------------

### GUEST LECTURES

Summer 2024	<b>Introduction to (Large) Language Models</b> , Host: Darin Gray	<b>USC Viterbi K-12 Discover Engineering</b>
Spring 2024	<b>Language Generation and Generative Evaluation</b> , Host: Xuezhe Ma	<b>USC CSCI 544: Advanced NLP</b>
Summer 2023	<b>Introduction to NLP</b> , Host: Darin Gray	<b>USC Viterbi K-12 Discover Engineering</b>
Spring 2023	<b>Contextualizing Bias</b> , Host: Maja Matarić	<b>USC CSCI 697: Seminar in CS Research</b>
Spring 2022	<b>What's in a dataset?</b> , Host: Cyrus Shahabi	<b>USC CSCI 697: Seminar in CS Research</b>
Spring 2022	<b>Data-Centric NLP</b> , Host: Dongyeop Kang	<b>UMinn. CSCI 8980: Intro / NLP Research</b>
Spring 2022	<b>Annotating Ambiguous NLI</b> , Host: Dirk Hovy	<b>MilaNLP: Coding Aperitivo</b>
Spring 2021	<b>Transfer Learning</b> , Host: Zaid Harachoui	<b>UW DATA 598: Statistical Deep Learning</b>
Fall 2020	<b>Biases and Interpretability in NLP</b> , Host: Eunsol Choi	<b>UT-Austin CS 395T: Topics in NLP</b>
Winter 2019	<b>Phrase-Structure Parsing</b> , Host: Noah A. Smith	<b>UW CSE 447/547M: NLP</b>
Spring 2018	<b>Minimum Bayes Risk Decoding</b> , Host: Waleed Ammar	<b>UW CSE 599 D1: Advanced Topics in NLP</b>
Spring 2018	<b>Dependency Parsing</b> , Host: Noah A. Smith	<b>UW CSEP 517: NLP</b>
Fall 2017	<b>Unsupervised Learning</b> , Host: Noah A. Smith	<b>UW CSE 446: Machine Learning</b>
Fall 2014	<b>Dependency Parsing with Chu-Liu-Edmonds</b> , Hosts: Chris Dyer, Alon Lavie, Bob Frederking	<b>CMU 11-271: Algorithms for NLP</b>