

Swabha Swayamdipta

GABILAN ASSISTANT PROFESSOR OF COMPUTER SCIENCE, UNIVERSITY OF SOUTHERN CALIFORNIA

✉ swabhas@usc.edu | 🏠 https://swabhs.com | 🌐 swabhs | 🐦 @swabhaz

Professional Experience

University of Southern California

GABILAN ASSISTANT PROFESSOR OF COMPUTER SCIENCE

- Associate Director of the Center for AI and Society, since Fall 2023

Los Angeles, CA, USA

Aug 2022 - PRESENT

Allen Institute for AI

POSTDOCTORAL INVESTIGATOR

- **Advisor:** Yejin Choi
- **Project Team:** MOSAIC

Seattle, WA, USA

Aug 2019 - June 2022

Oracle Server Technologies

MEMBER TECHNICAL STAFF

- **Project:** Unified messaging for Fusion Middleware

Bangalore, India

2010 - 2011

Education

Carnegie Mellon University

PHD IN LANGUAGE AND INFORMATION TECHNOLOGIES

- **Advisors:** Noah A. Smith, Chris Dyer
- **University of Washington**, Seattle, WA — *Visiting PhD student Fall 2015 - Summer 2019*
- **Thesis:** Syntactic Inductive Biases in NLP

Pittsburgh, PA, USA

2013 - 2019

Columbia University

MASTERS IN COMPUTER SCIENCE

- **Advisors:** Owen Rambow, Michael Collins
- **Specialization:** Natural Language Processing

City of New York, NY, USA

2011 - 2012

National Institute of Technology

BACHELORS OF TECHNOLOGY IN COMPUTER SCIENCE AND ENGINEERING

Calicut, Kerala, India

2006 - 2010

Awards and Grants

- | | | |
|------|---|------------------------------------|
| 2025 | Research Fellowship at the Special Program on LLMs and Transformers , Simons Institute | UC Berkeley |
| 2024 | CISE Core Medium RII: Hard Data to the Model , NSF | |
| 2024 | DEI in Research Fellowship , USC Zumberge | |
| 2023 | Rising Stars Award , Intel | |
| 2022 | Young Investigators Award , Allen Institute for AI | |
| 2022 | Gabilan Assistant Professor Fellowship , Women in Science and Engineering | USC |
| 2022 | Outstanding Paper Award , Understanding Dataset Difficulty with V-Usable Information | ICML |
| 2021 | Outstanding Paper Award , MAUVE: Measuring the Gap Between Neural Text and Human Text using Divergence Frontiers | NeurIPS |
| 2020 | Best Paper: Honorable Mention , Don't Stop Pretraining: Adapt Language Models to Domains and Tasks | ACL |
| 2014 | Best Student Presentation , Diversity in Dependency Parsing | CMU LTI Student Research Symposium |
| 2009 | Sun Campus Ambassador Scholarship , Bangalore, India | Sun Microsystems |
| 2003 | National Talent Search Scholarship , Rourkela, India | NCERT, India |

Advising

Fall 2024-Present	Gregory Yauney , PostDoc, USC	
Fall 2024-Present	Xinyue Cui , PhD, USC	
Fall 2024-Present	Muru Zhang , PhD, USC	<i>Co-Advisor: Robin Jia</i>
Fall 2024-Present	Atharva Kulkarni , PhD, USC	
Spring 2024-Present	Brihi Joshi , PhD, USC	<i>Co-Advisor: Xiang Ren</i>
Fall 2023-Present	Matthew Finlayson , PhD, USC	<i>Co-Advisor: Xiang Ren</i>
Fall 2022-Present	Sayan Ghosh , PhD, USC	
Fall 2023-Present	Jaspreet Ranjit , PhD, USC	
Spring 2024-Present	Shahzaib Saqib Warraich , Masters, USC	
Spring 2024-Present	Shauryasikt Jena , Masters, USC	
Spring 2024-Present	Xingjian Dong , Masters, USC	
Spring 2024-Present	Ryan Wang , BS, USC	<i>Co-Advisor: Robin Jia</i>
Spring 2024-Present	Joseph Liu , BS, USC	
Spring 2024-Present	Keyu He , BS, USC	
Fall 2023-Present	Risha Surana , BS, USC	
Summer 2023-Present	Aryan Gulati , BS, USC	<i>Co-Advisor: Antonio Ortega</i>
Spring 2023-2024	Xinyue Cui , Masters, USC	<i>Now PhD@USC</i>
Fall 2023-Spring 2024	Jacky Mo , BS, USC	
Summer-Fall 2023	Jayne Bottarini , BS, USC	
Summer-Fall 2023	Kritin Dhoka , BS, USC	
Summer-Fall 2023	Catherine He , BS, USC	
Spring 2023-2024	Yoonsoo Nam , Masters, USC	<i>Now at RenderWolf</i>
Spring-Summer 2023	Adam Lehavi , BS, USC	
Summer 2023	Ruyuan Zuo , Masters, USC	<i>Now at Google</i>
Spring 2022	Hanjie Chen / Junlin Wang / Abdallah Bashir , PhD, UVA / Masters, UCI / Masters, Saarland Universit	<i>AI2 Intern</i>
Fall 2021	Jillian Fisher / Liwei Jiang , PhD, UW	
Summer 2021	Kawin Ethayarajh , PhD, Stanford University	<i>AI2 Intern</i>
Summer 2021	Sarah Wiegrefe , PhD, George Institute of Technology	<i>AI2 Intern</i>
Summer 2021	Ximing Liu , BS, UW CSE	<i>AI2 Intern</i>
Fall 2020-2023	Alisa Liu , PhD, UW CSE	
Fall 2020	Alon Jacovi , PhD, Bar Ilan University	<i>AI2 Intern</i>
Summer 2020-1	Jenny Liang , BS, UW CSE	<i>AI2 Intern</i>
Summer 2020	Jize Cao , BS, UW CSE	<i>AI2 Intern</i>
Spring 2020-2023	Xuhui Zhou , CLMS, UW Linguistics	
Fall 2019 - Spring 2020	Yiben Yang , PhD, Northwestern University	
Fall 2019 - Spring 2020	Chaitanya Malaviya , PYI, Allen Institute for AI	
2018 - 2019	Ron Fan / Karishma Mandyam , MS, UW CSE / BS, UW CSE	

Publications

CONFERENCE PAPERS

Logits of API-Protected LLMs Leak Proprietary Information	<i>COLM 2024 (To Appear)</i>
• M. Finlayson, X. Ren, S. Swayamdipta ,	
Crowd-Calibrator: Can Annotator Disagreement Inform Calibration in Subjective Tasks?	<i>COLM 2024 (To Appear)</i>
• U. Khurana, E. Nalisnick, A. Fokkens, S. Swayamdipta ,	
Annotating FrameNet via Structure-Conditioned Language Generation	<i>ACL 2024 (To Appear)</i>
• X. Cui, and S. Swayamdipta	
Closing the Curious Case of Neural Text Degeneration	<i>ICLR 2024</i>
• M. Finlayson, J. Hewitt, A. Koller, S. Swayamdipta , A. Sabharwal	

Does Video Summarization Require Videos? Quantifying the Effectiveness of Language in Video Summarization	ICASSP 2024
• Y. Nam, A. Lehavi, D. Yang, D. Bose, S. Swayamdipta , and S. Narayanan	
NeuroComparatives: Neuro-Symbolic Distillation of Comparative Knowledge	NAACL-Findings 2024
• P. Howard, J. Wang, V. Lal, G. Singer, Y. Choi, and S. Swayamdipta	
Generative Explanations for Program Synthesizers	ICSE + VL/HCC 2024
• A. Nazari, S. Chattopadhyay, S. Swayamdipta , M. Raghothaman	
We're Afraid Language Models Aren't Modeling Ambiguity	EMNLP 2023
• A. Liu, Z. Wu, J. Michael, A. Suhr, P. West, A. Koller, S. Swayamdipta , N. A. Smith, and Y. Choi	
MAUVE Scores for Generative Models: Theory and Practice	JMLR 2023
• K. Pillutla, L. Liu, J. Thickstun, S. Welleck, S. Swayamdipta , R. Zellers, S. Oh, Y. Choi, and Z. Harchaoui	
I2D2: Inductive Knowledge Distillation with NeuroLogic and Self-Imitation	ACL 2023
• C. Bhagavatula, J. D. Hwang, D. Downey, R. Le Bras, X. Lu, K. Sakaguchi, S. Swayamdipta , P. West, and Y. Choi	
REV: Information-Theoretic Evaluation of Free-Text Rationales	ACL 2023
• H. Chen, F. Brahman, X. Ren, Y. Ji, Y. Choi, and S. Swayamdipta	
COBRA Frames: Contextual Reasoning about Effects and Harms of Offensive Statements	ACL-Findings 2023
• X. Zhou, H. Zhu, A. Yerukola, T. Davidson, J. D. Hwang, S. Swayamdipta , and M. Sap	
Investigating the Benefits of Free-Form Rationales	EMNLP-Findings 2022
J. Sun, S. Swayamdipta , J. May, and X. Ma	
NeuroCounterfactuals: Beyond Minimal-Edit Counterfactuals for Richer Data Augmentation	EMNLP-Findings 2022
P. Howard, G. Singer, V. Lal, Y. Choi, and S. Swayamdipta	
WaNLI: Worker and AI Collaboration for Natural Language Inference Dataset Creation	EMNLP-Findings 2022
A. Liu, S. Swayamdipta , N. A. Smith, and Y. Choi	
Reframing Human-AI Collaboration for Generating Free-Text Explanations	NAACL 2022
S. Wiegrefe, J. Hessel, S. Swayamdipta , M. Riedl, and Y. Choi	
Annotators with Attitudes: How Annotator Beliefs And Identities Bias Toxic Language Detection	NAACL 2022
M. Sap, S. Swayamdipta , L. Vianna, X. Zhou, Y. Choi, and N. A. Smith	
Understanding Dataset Difficulty with V-Usable Information	ICML 2022
• K. Ethayarajh, Y. Choi, and S. Swayamdipta	
• Outstanding Paper Award	
MAUVE: Measuring the Gap Between Neural Text and Human Text using Divergence Frontiers	NeurIPS 2021
• K. Pillutla, S. Swayamdipta , R. Zellers, J. Thickstun, S. Wellecks, Y. Choi, and Z. Harchaoui	
• Outstanding Paper Award	
Contrastive Explanations for Model Interpretability	EMNLP 2021
A. Jacovi, S. Swayamdipta , S. Ravfogel, Y. Elazar, Y. Choi, and Y. Goldberg	
On-the-Fly Controlled Text Generation with Experts and Anti-Experts	ACL 2021
A. Liu, M. Sap, X. Lu, S. Swayamdipta , C. Bhagavatula, N. A. Smith, and Y. Choi	
Challenges in Automated Debiasing for Toxic Language Detection	EACL 2021
• X. Zhou, M. Sap, S. Swayamdipta , N. A. Smith, and Y. Choi	
Dataset Cartography: Mapping and Diagnosing Datasets with Training Dynamics	EMNLP 2020
• S. Swayamdipta , R. Schwartz, N. Lourie, Y. Wang, H. Hajishirzi, N. A. Smith, Y. Choi	
Generative Data Augmentation for Commonsense Reasoning	EMNLP-Findings 2020
• Y. Yang, C. Malaviya, J. Fernandez, S. Swayamdipta , R. LeBras, J. Wang, C. Bhagavatula, Y. Choi, and D. Downey	
Adversarial Filters of Dataset Biases	ICML 2020
• R. LeBras, S. Swayamdipta , C. Bhagavatula, R. Zellers, M. E. Peters, A. Sabharwal, and Y. Choi	
The Right Tool for the Job: Matching Model and Instance Complexities	ACL 2020
• R. Schwartz, G. Stanovsky, S. Swayamdipta , J. Dodge, and N. A. Smith	

Don't Stop Pretraining: Adapt Language Models to Domains and Tasks	ACL 2020
<ul style="list-style-type: none"> S. Gururangan, A. Marasović, S. Swayamdipta, K. Lo, I. Beltagy, D. Downey, and N. A. Smith Best Paper Honorable Mention 	
Syntactic Scaffolds for Semantic Structures	EMNLP 2018
<ul style="list-style-type: none"> S. Swayamdipta, S. Thomson, K. Lee, L. Zettlemoyer, C. Dyer, N. A. Smith. 	
Learning Joint Semantic Parsers from Disjoint Data	NAACL 2018
<ul style="list-style-type: none"> H. Peng, S. Thomson, S. Swayamdipta, N. A. Smith 	
Annotation Artifacts in Natural Language Inference Data	NAACL 2018
<ul style="list-style-type: none"> S. Gururangan*, S. Swayamdipta*, O. Levy, R. Schwartz, S. Bowman, and N. A. Smith * equal contribution 	
Polyglot Semantic Role Labeling	ACL 2018
<ul style="list-style-type: none"> P. Mulcaire, S. Swayamdipta, and N. A. Smith 	
Multi-Mention Learning for Reading Comprehension with Neural Cascades	ICLR 2018
<ul style="list-style-type: none"> S. Swayamdipta, A. Parikh, T. Kwiatkowski 	
Greedy, Joint Syntactic and Semantic Parsing with Stack LSTMs	CoNLL 2016
<ul style="list-style-type: none"> S. Swayamdipta, M. Ballesteros, C. Dyer, N. A. Smith 	
A Dependency Parser for Tweets	EMNLP 2014
<ul style="list-style-type: none"> L. Kong, N. Schneider, S. Swayamdipta, A. Bhatia, C. Dyer, N. A. Smith 	
The Pursuit of Power and its Manifestation in Written Dialog	ICSC 2012
<ul style="list-style-type: none"> S. Swayamdipta, O. Rambow 	

WORKING PAPERS

Compare without Despair: Reliable Preference Evaluation with Generation Separability	Under Submission
<ul style="list-style-type: none"> S. Ghosh, T. Srinivasan, S. Swayamdipta, 	
Out-of-Distribution Detection through Soft Clustering with Non-Negative Kernel Regression	Under Submission
<ul style="list-style-type: none"> A. Gulati, X. Dong, C. Hurtado, S. Shekizhar, S. Swayamdipta, A. Ortega 	
OATH-Frames: Characterizing Online Attitudes towards Homelessness via LLM Assistants	Under Submission
<ul style="list-style-type: none"> J. Ranjit, B. Joshi, R. Dorn, L. Petry, O. Koumoundouros, J. Bottarini, P. Liu, E. Rice, and S. Swayamdipta Best Poster @ ShowCAIS 2024 	
Shallow Syntax in Deep Water	arXiv:1908.11047
<ul style="list-style-type: none"> S. Swayamdipta, M. Peters, B. Roof, C. Dyer and N. A. Smith 	
Frame-Semantic Parsing with Softmax-Margin Segmental RNNs and a Syntactic Scaffold	arXiv:1706.09528
<ul style="list-style-type: none"> S. Swayamdipta, S. Thomson, C. Dyer and N. A. Smith 	
DyNet: The Dynamic Neural Network Toolkit	arXiv:1701.03980
<ul style="list-style-type: none"> G. Neubig, C. Dyer, Y. Goldberg, A. Matthews, W. Ammar, A. Anastasopoulos, M. Ballesteros, D. Chiang, D. Clothiaux, T. Cohn, K. Duh, M. Faruqui, C. Gan, D. Garrette, Y. Ji, L. Kong, A. Kuncoro, G. Kumar, C. Malaviya, P. Michel, Y. Oda, M. Richardson, N. Saphra, S. Swayamdipta, P. Yin 	

WORKSHOP PAPERS

Sister Help: Data Augmentation for Frame-Semantic Role Labeling	LAW-DMR @ EMNLP 2021
<ul style="list-style-type: none"> A. Pancholy, S. Swayamdipta, M. R. L. Petruck 	
Multi-Task Learning for Incremental Parsing using Stack LSTMs	WiML @ NeurIPS 2016
<ul style="list-style-type: none"> S. Swayamdipta, M. Ballesteros, C. Dyer, N. A. Smith 	
CMU: Arc-Factored, Discriminative Semantic Dependency Parsing	SemEval 2014
<ul style="list-style-type: none"> S. Thomson, D. Bamman, J. Dodge, S. Swayamdipta, N. Schneider, C. Dyer, N. A. Smith 	
The CMU Machine Translation Systems	WMT 2014
<ul style="list-style-type: none"> A. Matthews, C. Dyer, A. Lavie, G. Hanneman, W. Ammar, A. Bhatia, S. Swayamdipta, E. Schlinger, Y. Tsvetkov 	

Invited Talks

Dec 5, 2024	(Upcoming) Learning Machine Seminar Series (LMSS),	Cornell Tech
Oct 10, 2024	(Upcoming) NLP/Text-as-Data Talk Series,	New York University
May 3, 2024	Understanding LLMs through their Generative Behavior, Successes and Shortcomings, ISI AI Seminar	USC ISI
Mar 26, 2024	Towards (Closed-Source) LLM Accountability via Logit Signatures, NSF-Open Source Generative AI Workshop	Cornell Tech
Mar 13, 2024	Understanding LLMs through their Generative Behavior, Successes and Shortcomings, Data Science Seminar	University of Utah
Feb 13, 2024	LLM Inference for Scaling Data Creation, Intel Rising Stars Talk	Intel Labs
Feb 8, 2024	Understanding LLMs via their Generative Successes and Shortcomings, LTL Seminar	University of Cambridge
Dec 15, 2023	Understanding LLMs via their Generative Successes and Shortcomings, Keynote	ATTRIB Workshop @ NeurIPS 2023
Nov 16, 2023	Understanding Online Discourse through Social Context and Structured Pragmatics, UCLA Communications Symposium	UCLA
Oct 3, 2023	Understanding Data with \mathcal{V}-Information, Keynote	Towards the Next Generation of Computer Vision Datasets Workshop @ ICCV 2023
Sep 22, 2023	The Role of Language Models in Natural Language Processing, Los Angeles County of Education Computer Science Speaker Series	Los Angeles
Aug 24, 2023	Understanding Datasets and Explanations through \mathcal{V}-Information, Alumni Keynote: CMU LTI Student Research Symposium	CMU, Pittsburgh
Jul 13, 2023	Contextualizing Representations in Varied Annotator Perspectives, Rep4NLP Workshop	ACL Toronto
Jun 14, 2023	Contextualizing Data in Varied Annotator Perspectives, Responsible Machine Learning	Google Brain
Mar 14, 2023	Understanding Dataset Difficulty with \mathcal{V}-Usable Information, Cohere for AI Reading Group	Cohere for AI
Feb 28, 2023	Designing Controls and Filters for Dataset Generation, Spotify Research Seminar	Spotify Research Labs
Feb 03, 2023	What's in a Dataset? Interpreting datasets to enable better data creation, Amazon Data-Centric AI Seminar	Amazon
Feb 01, 2023	Contextualizing Bias in Hate Speech Detection, CAIS++ Seminar	USC
Nov 18, 2022	Generating Datasets for Robust Generalization, SoCal NLP Symposium	UC Santa Barbara
Nov 09, 2022	Contextualizing Bias in Hate Speech Detection through Annotator Perspectives, Center for AI in Society (CAIS) Seminar	USC
May 23, 2022	The Devil's in the Data: Mapping and Generating Datasets for Robust Generalization, ACL Spotlight Talks for Young Rising Stars	Dublin, Ireland
May 13, 2022	Mapping and Generating Datasets for Robust Generalization, CS Seminar	UC Irvine
Mar 16, 2022	Rethinking Dataset Construction: : The Role of Generative Modeling and Annotator Perspectives, IFDS Ethics Seminar	UC Santa Cruz
Oct 20, 2021	What's in your Data? Mapping Datasets and Exploring Data Usability, Machine Learning Research Group Talks	Oracle
Feb 24, 2020	Addressing Biases for Robust, Generalizable AI, NERT Seminar	Georgetown University
Feb 12, 2020	Addressing Biases for Robust, Generalizable AI, NLP Seminar	Georgia Tech
Nov 02, 2020	Responsible AI: Addressing Biases in Datasets and Models, E+D Product Leads	Microsoft
Oct 30, 2019	Biases and Learning Challenges in Natural Language Processing, Rising Stars in EECS	UIUC
Oct 18, 2019	Sprucing up a Dataset: Adversarially Filtering Dataset Artifacts, Linguistics Colloquium	UW Linguistics
Apr 08, 2019	Learning Challenges in Natural Language Processing, AI2 Seminar	Allen Institute for AI
Apr 04, 2019	Learning Challenges in Natural Language Processing, NLP Seminar	Stanford University
Mar 13, 2019	Learning Challenges in Natural Language Processing, CLSP Seminar	Johns Hopkins University
Mar 11, 2019	Learning Challenges in Natural Language Processing, CILVR Seminar	New York University
Jan 22, 2019	Learning Challenges in Natural Language Processing, MSR Seminar	Microsoft Research, New York
Sep 15, 2018	Representation Learning with Linguistic Structure, MSR AI Breakthroughs Workshop	Microsoft Research, Redmond
Apr 27, 2018	Syntactic Scaffolds for Semantic Structures, NorthWest NLP	Microsoft Research, Redmond

Invited Panels

Oct 18, 2024 **(Upcoming) Alfred P. Sloan Foundation Science Seminar,**
 Jul 13, 2023 **Limitations of Large Language Models,** Rep4NLP Workshop
 Jun 29, 2023 **Generative AI and Education,** Roundtable
 Apr 5, 2023 **Almaginings,** Polymathic Pizza
 Jul 14, 2022 **Adversarial Data Augmentations,** DADC Workshop
 Jun 17, 2022 **Role of LLMs,** Responsible AI Symposium

USC School of Cinematic Arts
ACL Toronto
ASEE 2023
USC Sidney Harmon Academy
for Polymathic Studies
NAACL Seattle
AILA

Professional Service

EXTERNAL SERVICE

Senior Area Chair

- EMNLP 2024: Machine Learning in NLP
- ACL 2024: Sentence-level Semantics

Area Chair

- ICLR 2025 (upcoming)
- COLM 2024
- ICLR 2024
- EMNLP 2023: Machine Learning for NLP
- ACL 2023: Interpretability
- EMNLP 2022: Language Models
- EMNLP 2021: Machine Learning for NLP
- NAACL 2021: Sentence-Level Semantics
- EACL 2021: Sentence-Level Semantics
- ACL 2020: Semantics (Long)

Reviewer: Conferences

- ACL (ARR) 2015-2022
- NAACL 2015-2022
- EMNLP 2015-2022
- NeurIPS 2018-2022
- ICML 2015, 2019, 2020, 2023
- EACL 2017
- AAAI 2017-2020
- CoNLL 2017-2018, 2020

Reviewer: Workshops

- NAACL Student Research Workshop 2016
- Workshop for Women in Machine Learning (WiML) at NeurIPS 2016
- Machine Reading for Question Answering (MRQA) Workshop 2017

Reviewer: Journals

- Transactions of ACL 2020 - PRESENT
- Computational Linguistics 2019 - 2022
- Journal of AI Research 2019

Co-organizer

- EACL 2024: Workshop on Uncertainty in NLP 2024
- NAACL 2022: Workshop on Deep Learning in Low Resource NLP (DeepLo) 2022
- EMNLP 2019: Workshop on Deep Learning in Low Resource NLP (DeepLo) 2019
- West Coast NLP Workshop 2018
- Women's Research Day 2017
- UW-NLP Retreat 2015-2016

Mentorship at Conferences

- EMNLP 2020, ACL 2020, EMNLP 2018

INTERNAL SERVICE

Univeristy of Southern California

- Diversity, Equity and Inclusion Committee
- Apple and IBM Fellowship Committee
- Distinguished Lectures Committee
- Annual Faculty Review Committee
- Apple and IBM Fellowship Committee
- PhD Fellowship Admissions Committee
- Viterbi Undergraduate Admission Scholarship

Los Angeles, CA, USA

2023-present
2024
2023-2024
2023-2024
2023
2022-2023
2023-2024

Allen Institute for AI

- Diversity, Equity and Inclusion Committee
- AllenAI Outstanding Engineer Scholarship Committee

Seattle WA, USA

2019-2022
2019

Paul G. Allen School of CSE, University of Washington

- Faculty Recruiting Liaison
- Graduate Student Advisory Council
- Outreach: Highline Public High School
- CS: Gradswomen

Seattle WA, USA

2019
2018
2017
2016-2019

Teaching

CLASSES

Fall 2024	CSCI 544: Applied NLP , Students: 260	USC
Spring 2024	CSCI 499: Language Models in NLP , Students: 28	USC
Fall 2023	CSCI 499: Language Models in NLP , Students: 30	USC
Fall 2022	CSCI 699: Data-Centric NLP , Students: 29	USC

TUTORIALS

2022	OAMLS: Online Asian Machine Learning School , Attendees: 350+	ACML
2019	Transfer Learning in Natural Language Processing , Attendees: 350+	NAACL
2018	Frame Semantics across Languages: Towards a Multilingual FrameNet , Attendees: 30+	CoLING

TEACHING ASSISTANTSHIP

Spring 2015	CMU 11-411/611: Natural Language Processing , Instructors: Chris Dyer, Alan W Black, Shomir Wilson	Head TA
-------------	---	----------------

GUEST LECTURES

Summer 2024	Introduction to (Large) Language Models , Host: Darin Gray	USC Viterbi K-12 Discover Engineering
Spring 2024	Language Generation and Generative Evaluation , Host: Xuezhe Ma	USC CSCI 544: Advanced NLP
Summer 2023	Introduction to NLP , Host: Darin Gray	USC Viterbi K-12 Discover Engineering
Spring 2023	Contextualizing Bias , Host: Maja Matarić	USC CSCI 697: Seminar in CS Research
Spring 2022	What's in a dataset? , Host: Cyrus Shahabi	USC CSCI 697: Seminar in CS Research
Spring 2022	Data-Centric NLP , Host: Dongyeop Kang	UMinn. CSCI 8980: Intro / NLP Research
Spring 2022	Annotating Ambiguous NLI , Host: Dirk Hovy	MilaNLP: Coding Aperitivo
Spring 2021	Transfer Learning , Host: Zaid Harachoui	UW DATA 598: Statistical Deep Learning
Fall 2020	Biases and Interpretability in NLP , Host: Eunsol Choi	UT-Austin CS 395T: Topics in NLP
Winter 2019	Phrase-Structure Parsing , Host: Noah A. Smith	UW CSE 447/547M: NLP
Spring 2018	Minimum Bayes Risk Decoding , Host: Waleed Ammar	UW CSE 599 D1: Advanced Topics in NLP
Spring 2018	Dependency Parsing , Host: Noah A. Smith	UW CSEP 517: NLP
Fall 2017	Unsupervised Learning , Host: Noah A. Smith	UW CSE 446: Machine Learning
Fall 2014	Dependency Parsing with Chu-Liu-Edmonds , Hosts: Chris Dyer, Alon Lavie, Bob Frederking	CMU 11-271: Algorithms for NLP